

Final Project Code

April 10, 2024

```
[ ]: import pandas as pd
import numpy as np
import copy
```

```
[ ]: dataset = pd.read_csv('./dataset_ideal.csv')
```

```
[ ]: # Data Exploration
dataset.shape[0]
dataset.head()
```

```
[ ]:      age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope  \
0    52    1   0      125    212   0         1      168     0         1.0     2
1    53    1   0      140    203   1         0      155     1         3.1     0
2    70    1   0      145    174   0         1      125     1         2.6     0
3    61    1   0      148    203   0         1      161     0         0.0     2
4    62    0   0      138    294   1         1      106     0         1.9     1
```

```
      ca  thal  target
0     2     3         0
1     0     3         0
2     0     3         0
3     1     3         0
4     3     2         0
```

```
[ ]: dataset.tail()
```

```
[ ]:      age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  \
1020   59    1   1      140    221   0         1      164     1         0.0
1021   60    1   0      125    258   0         0      141     1         2.8
1022   47    1   0      110    275   0         0      118     1         1.0
1023   50    0   0      110    254   0         0      159     0         0.0
1024   54    1   0      120    188   0         1      113     0         1.4
```

```
      slope  ca  thal  target
1020     2    0     2         1
1021     1    1     3         0
1022     1    1     2         0
1023     2    0     2         1
```

```
1024      1      1      3      0
```

```
[ ]: dataset.describe()
```

```
[ ]:
```

	age	sex	cp	trestbps	chol \
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
mean	54.434146	0.695610	0.942439	131.611707	246.000000
std	9.072290	0.460373	1.029641	17.516718	51.59251
min	29.000000	0.000000	0.000000	94.000000	126.000000
25%	48.000000	0.000000	0.000000	120.000000	211.000000
50%	56.000000	1.000000	1.000000	130.000000	240.000000
75%	61.000000	1.000000	2.000000	140.000000	275.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000

	fbs	restecg	thalach	exang	oldpeak \
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
mean	0.149268	0.529756	149.114146	0.336585	1.071512
std	0.356527	0.527878	23.005724	0.472772	1.175053
min	0.000000	0.000000	71.000000	0.000000	0.000000
25%	0.000000	0.000000	132.000000	0.000000	0.000000
50%	0.000000	1.000000	152.000000	0.000000	0.800000
75%	0.000000	1.000000	166.000000	1.000000	1.800000
max	1.000000	2.000000	202.000000	1.000000	6.200000

	slope	ca	thal	target
count	1025.000000	1025.000000	1025.000000	1025.000000
mean	1.385366	0.754146	2.323902	0.513171
std	0.617755	1.030798	0.620660	0.500070
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	2.000000	0.000000
50%	1.000000	0.000000	2.000000	1.000000
75%	2.000000	1.000000	3.000000	1.000000
max	2.000000	4.000000	3.000000	1.000000

```
[ ]: dataset.isnull().sum()
```

```
[ ]: age      0
      sex      0
      cp      0
      trestbps 0
      chol     0
      fbs      0
      restecg  0
      thalach  0
      exang    0
      oldpeak  0
      slope    0
```

```
ca          0
thal        0
target      0
dtype: int64
```

```
[ ]: column_dtypes = dataset.dtypes

print("Data types of each column:")
print(column_dtypes)
```

Data types of each column:

```
age          int64
sex          int64
cp           int64
trestbps     int64
chol         int64
fbs          int64
restecg      int64
thalach      int64
exang        int64
oldpeak      float64
slope        int64
ca           int64
thal         int64
target       int64
dtype: object
```

```
[ ]: %matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[ ]: import pandas_profiling as pp
pp.ProfileReport(dataset)
```

Summarize dataset: 0%| | 0/5 [00:00<?, ?it/s]

Generate report structure: 0%| | 0/1 [00:00<?, ?it/s]

Render HTML: 0%| | 0/1 [00:00<?, ?it/s]

<IPython.core.display.HTML object>

```
[ ]:
```

```
[ ]: class_distribution = dataset['target'].value_counts()
# Labels
labels = class_distribution.index.astype(str)
targets = ['no disease', 'disease']

# Plot
```


```

fig, ax = plt.subplots(figsize=(6, 4))
x = range(len(labels))
counts = class_distribution.values
bars = ax.bar(x, counts, color=['navy', 'orange'])
plt.title('Class Distribution')
total_count_0 = class_distribution[0]
total_count_1 = class_distribution[1]
plt.text(0.5, -100, f'Total 0s: {total_count_0} Total 1s: {total_count_1}', u
        ha='center', fontsize=12)

plt.ylabel('Count')
plt.xticks(class_distribution.keys())

plt.show()

```

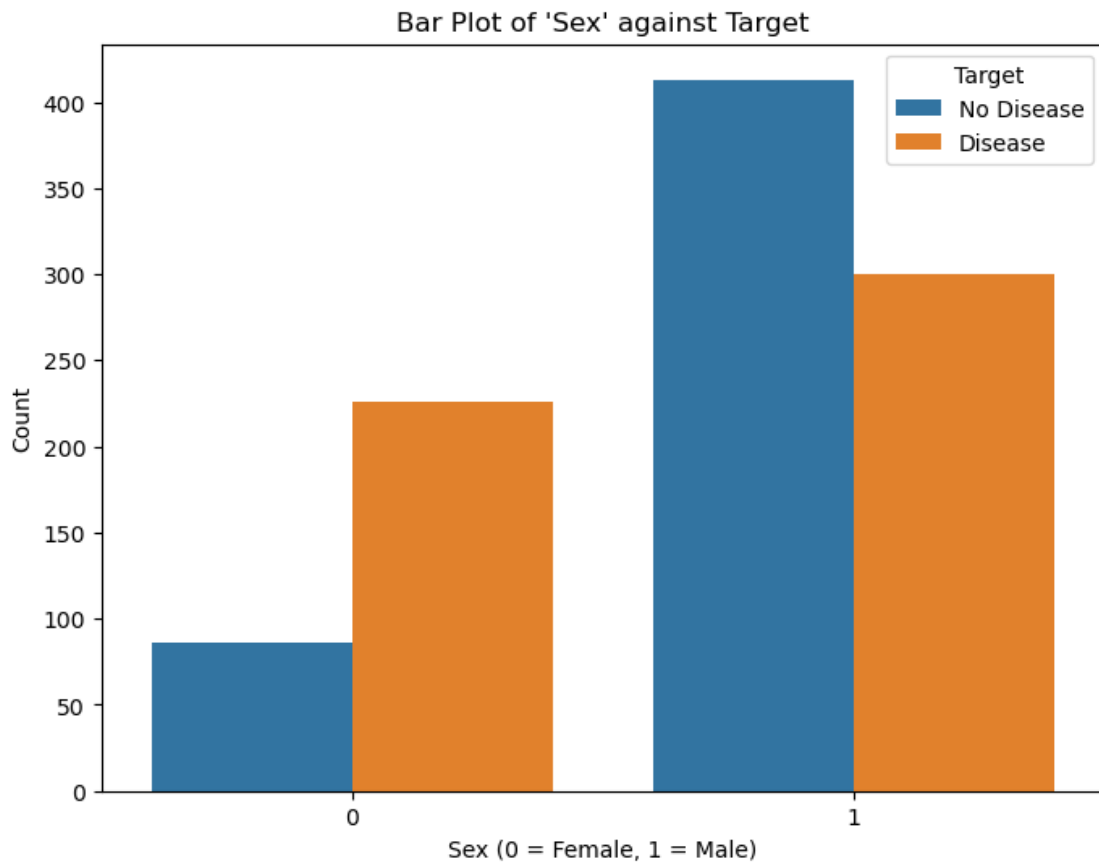


```

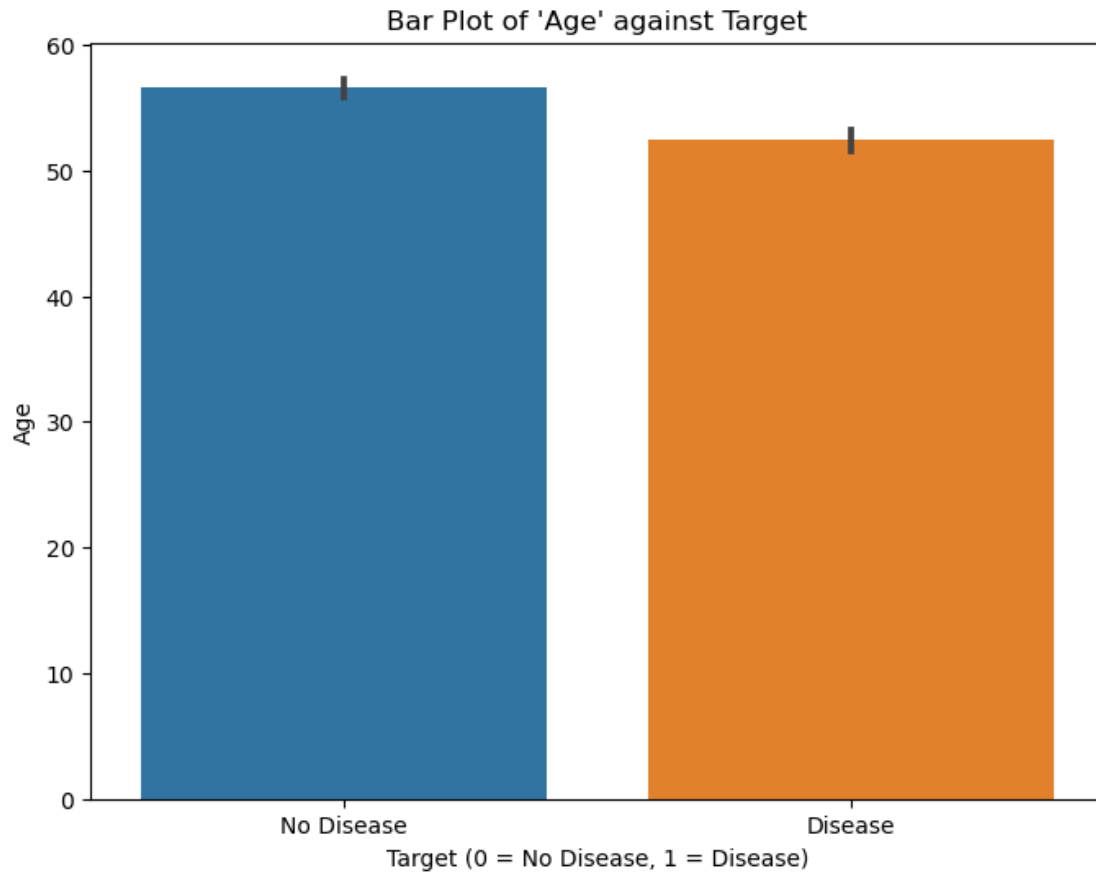
[ ]: # Bar plot for 'sex' against the target
plt.figure(figsize=(8, 6))
sns.countplot(x='sex', hue='target', data=dataset)
plt.title("Bar Plot of 'Sex' against Target")
plt.xlabel("Sex (0 = Female, 1 = Male)")
plt.ylabel("Count")
plt.legend(title='Target', labels=['No Disease', 'Disease'])

```


```
plt.show()
```



```
[ ]: # Box plot for 'age' against the target
plt.figure(figsize=(8, 6))
sns.barplot(x='target', y='age', data=dataset)
plt.title("Bar Plot of 'Age' against Target")
plt.xlabel("Target (0 = No Disease, 1 = Disease)")
plt.ylabel("Age")
plt.xticks(ticks=[0, 1], labels=['No Disease', 'Disease'])
plt.show()
```



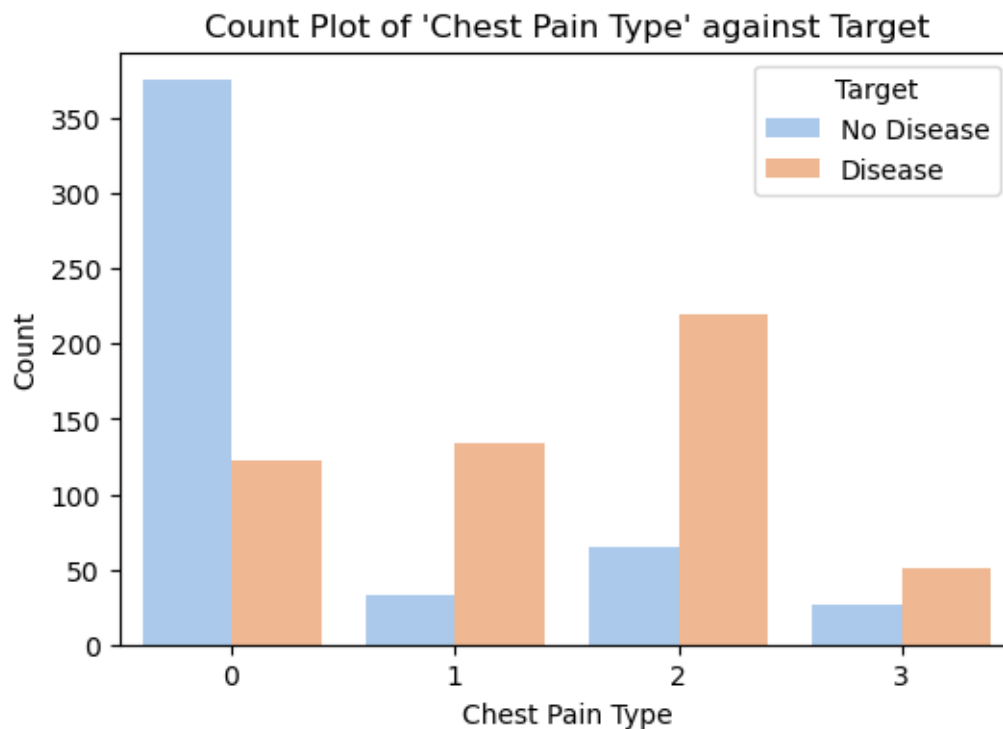
```
[ ]: # Box plot for 'cholesterol' against the target
plt.figure(figsize=(6, 4))
sns.barplot(x='target', y='chol', data=dataset, palette='coolwarm')
plt.title("Cholesterol Levels by Heart Disease Presence")
plt.xlabel("Presence of Heart Disease (0 = No, 1 = Yes)")
plt.ylabel("Cholesterol Level")
plt.xticks(ticks=[0, 1], labels=['No Heart Disease', 'Heart Disease'])
plt.show()
```



```
[ ]: # Box plot for 'chest pain type' against the target
plt.figure(figsize=(6, 4))
sns.countplot(x='cp', hue='target', data=dataset, palette='pastel')
plt.title("Count Plot of 'Chest Pain Type' against Target")
plt.xlabel("Chest Pain Type")
plt.ylabel("Count")
plt.legend(title='Target', labels=['No Disease', 'Disease'])
plt.show()

# Count of chest pain type by target
chest_pain_counts = dataset.groupby(['cp', 'target']).size().
    ↪reset_index(name='Count')

# Print the count information
print("Count of Chest Pain Type by Target:")
print(chest_pain_counts)
```



Count of Chest Pain Type by Target:

cp	target	Count
0	0	375
1	0	122
2	1	33
3	1	134
4	2	65
5	2	219
6	3	26
7	3	51


[]: *#Count Plot of Exercise Induced Angina*

```
plt.figure(figsize=(6, 4))
sns.countplot(x='exang', hue='target', data=dataset, palette='Set3')
plt.title("Exercise-Induced Angina by Heart Disease Presence")
plt.xlabel("Exercise-Induced Angina (0 = No, 1 = Yes)")
plt.ylabel("Count")
plt.legend(title="Presence of Heart Disease", labels=['No', 'Yes'])
plt.show()

exeang_counts = dataset.groupby(['exang', 'target']).size().
    ↪reset_index(name='Count')
```



```
# Print the count information
print("Count of Exercise-Induced Angina by Target:")
print(chest_pain_counts)
```



Count of Exercise-Induced Angina by Target:

	cp	target	Count
0	0	0	375
1	0	1	122
2	1	0	33
3	1	1	134
4	2	0	65
5	2	1	219
6	3	0	26
7	3	1	51

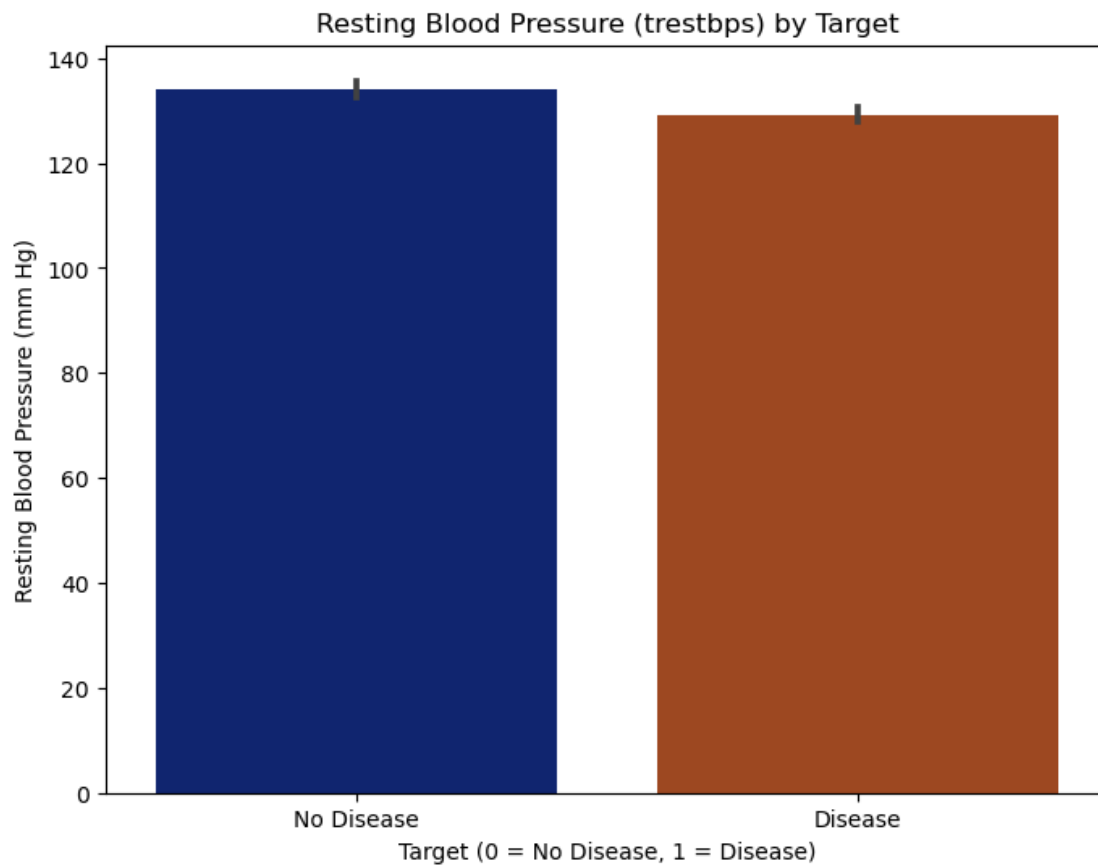
```
[ ]: plt.figure(figsize=(8, 6))
sns.barplot(x='target', y='trestbps', data=dataset, palette='dark')
plt.title("Resting Blood Pressure (trestbps) by Target")
plt.xlabel("Target (0 = No Disease, 1 = Disease)")
plt.ylabel("Resting Blood Pressure (mm Hg)")
plt.xticks(ticks=[0, 1], labels=['No Disease', 'Disease'])
plt.show()
```

```

exeang_counts = dataset.groupby(['trestbps', 'target']).size().
    ↪reset_index(name='Count')

# Print the count information
print("Count of Resting Blood Pressure by Target:")
print(chest_pain_counts)

```




Count of Resting Blood Pressure by Target:

	cp	target	Count
0	0	0	375
1	0	1	122
2	1	0	33
3	1	1	134
4	2	0	65
5	2	1	219
6	3	0	26
7	3	1	51

```
[ ]: # Compute the correlation matrix
correlation_matrix = dataset.corr()

# Plot the heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f",
            ↳annot_kws={"size": 10})
plt.title("Correlation Heatmap of All Features")
plt.show()
```



```
[ ]: # cp (chest pain type) - Strongest positive correlation (0.43)
# thalach (maximum heart rate achieved) - Second strongest positive correlation
↳ (0.42)
# exang (exercise induced angina) - Strongest negative correlation (-0.44)
# oldpeak (ST depression induced by exercise) - Strongest negative correlation
↳ (-0.44)
```

```

# trestbps (resting blood pressure) - Second strongest positive correlation (0.
↳42)
# ca (number of major vessels colored) - Second strongest negative correlation
↳(-0.38)
# thal (thalassemia) - Third strongest negative correlation (-0.34)
# slope (the slope of the peak exercise ST segment) - Third strongest negative
↳correlation (0.34)

```

```

[ ]: # Check number of unique values for each categorical feature
categorical_columns = []
for column in dataset.columns:
    if dataset[column].dtype == 'object' or dataset[column].nunique() < 10:
        categorical_columns.append(column)

print("Categorical columns:", categorical_columns)
for feature in categorical_columns:
    unique_values = dataset[feature].nunique()
    print(f"Feature '{feature}' has {unique_values} unique values:
↳{dataset[feature].unique()}")

```

```

Categorical columns: ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'ca',
'thal', 'target']
Feature 'sex' has 2 unique values: [1 0]
Feature 'cp' has 4 unique values: [0 1 2 3]
Feature 'fbs' has 2 unique values: [0 1]
Feature 'restecg' has 3 unique values: [1 0 2]
Feature 'exang' has 2 unique values: [0 1]
Feature 'slope' has 3 unique values: [2 0 1]
Feature 'ca' has 5 unique values: [2 0 1 3 4]
Feature 'thal' has 4 unique values: [3 2 1 0]
Feature 'target' has 2 unique values: [0 1]

```

```

[ ]: X = dataset.drop('target',axis=1)
y = dataset["target"]

# data splitting
from sklearn.model_selection import train_test_split

# Assuming X and y are your feature matrix and target vector
X_train_val, X_test, y_train_val, y_test = train_test_split(X, y, test_size=0.
↳20, random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X_train_val, y_train_val,
↳test_size=0.25, random_state=42)

# Print the dimensions of each set
print("Training set:", X_train.shape, y_train.shape)
print("Validation set:", X_val.shape, y_val.shape)

```

```
print("Test set:", X_test.shape, y_test.shape)
```

Training set: (615, 13) (615,)
Validation set: (205, 13) (205,)
Test set: (205, 13) (205,)

```
[ ]: #Data Preprocessing  
from sklearn.preprocessing import StandardScaler  
  
scaler = StandardScaler()  
X_train = scaler.fit_transform(X_train)  
X_val = scaler.transform(X_val)  
X_test = scaler.transform(X_test)
```

```
[ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score,  
     ↪ confusion_matrix, roc_auc_score, roc_curve, auc
```

```
[ ]: #Logistic Regression  
  
from sklearn.linear_model import LogisticRegression  
logistic_regression = LogisticRegression(max_iter=10000)  
  
loss_values = []  
# for i in range(1, 5001, 50): # Assuming 100 iterations  
#     loss = -np.sum(np.log(logistic_regression.predict_proba(X_train)[np.  
     ↪ arange(len(X_train)), y_train])) / len(X_train)  
#     print(f"epoch {i}, loss {loss}")  
#     loss_values.append(loss)  
  
logistic_regression.fit(X_train, y_train)
```

```
[ ]: LogisticRegression(max_iter=10000)
```

```
[ ]: # Logistic Regression: Make predictions on the test set  
y_pred_lr = logistic_regression.predict(X_test)  
  
# Calculate evaluation metrics  
accuracy = accuracy_score(y_test, y_pred_lr)  
precision = precision_score(y_test, y_pred_lr)  
recall = recall_score(y_test, y_pred_lr)  
conf_matrix = confusion_matrix(y_test, y_pred_lr)  
tn, fp, fn, tp = conf_matrix.ravel()  
specificity = tn / (tn + fp)  
f1_score = 2 * (precision * recall) / (precision + recall)  
roc_auc = roc_auc_score(y_test, y_pred_lr)  
  
# Print evaluation metrics  
print("Accuracy:", accuracy)
```


```

print("Precision:", precision)
print("Recall:", recall)
print("Specificity:", specificity)
print("F1 Score:", f1_score)
print("ROC AUC Score:", roc_auc)
print(conf_matrix)

plt.figure(figsize=(3, 3))
sns.heatmap(conf_matrix, annot=True, cmap='Blues', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

```

Accuracy: 0.7951219512195122
 Precision: 0.7520661157024794
 Recall: 0.883495145631068
 Specificity: 0.7058823529411765
 F1 Score: 0.8125000000000001
 ROC AUC Score: 0.7946887492861221
 [[72 30]
 [12 91]]



```

[ ]: #Naive Bayes

from sklearn.naive_bayes import GaussianNB

```

```
naive_bayes_classifier = GaussianNB()

naive_bayes_classifier.fit(X_train, y_train)
```

```
[ ]: GaussianNB()
```


```
[ ]: # Naive Bayes: Make predictions on the test set
y_pred_nb = naive_bayes_classifier.predict(X_test)

# Calculate evaluation metrics
accuracy_nb = accuracy_score(y_test, y_pred_nb)
precision_nb = precision_score(y_test, y_pred_nb)
recall_nb = recall_score(y_test, y_pred_nb)
conf_matrix_nb = confusion_matrix(y_test, y_pred_nb)
tn, fp, fn, tp = conf_matrix_nb.ravel()
specificity_nb = tn / (tn + fp)
f1_score_nb = 2 * (precision_nb * recall_nb) / (precision_nb + recall_nb)
roc_auc_nb = roc_auc_score(y_test, y_pred_nb)

# Print evaluation metrics
print("Accuracy:", accuracy_nb)
print("Precision:", precision_nb)
print("Recall:", recall_nb)
print("Specificity:", specificity_nb)
print("F1 Score:", f1_score_nb)
print("ROC AUC Score:", roc_auc_nb)

print(conf_matrix_nb)
plt.figure(figsize=(3, 3))
sns.heatmap(conf_matrix_nb, annot=True, cmap='Oranges', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```

```
Accuracy: 0.8048780487804879
Precision: 0.7560975609756098
Recall: 0.9029126213592233
Specificity: 0.7058823529411765
F1 Score: 0.8230088495575221
ROC AUC Score: 0.8043974871501998
[[72 30]
 [10 93]]
```



```
[ ]: #Random Forest

from sklearn.ensemble import RandomForestClassifier

random_forest_classifier = RandomForestClassifier()

random_forest_classifier.fit(X_train, y_train)
```

```
[ ]: RandomForestClassifier()
```

```
[ ]: # Random Forest: Make predictions on the test set
y_pred_rf = random_forest_classifier.predict(X_test)

# Calculate evaluation metrics
accuracy_rf = accuracy_score(y_test, y_pred_rf)
precision_rf = precision_score(y_test, y_pred_rf)
recall_rf = recall_score(y_test, y_pred_rf)
conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)
specificity_rf = tn / (tn + fp)
f1_score_rf = 2 * (precision_rf * recall_rf) / (precision_rf + recall_rf)
roc_auc_rf = roc_auc_score(y_test, y_pred_rf)

# Print evaluation metrics
print("Accuracy:", accuracy_rf)
print("Precision:", precision_rf)
print("Recall:", recall_rf)
print("Specificity:", specificity_rf)
```



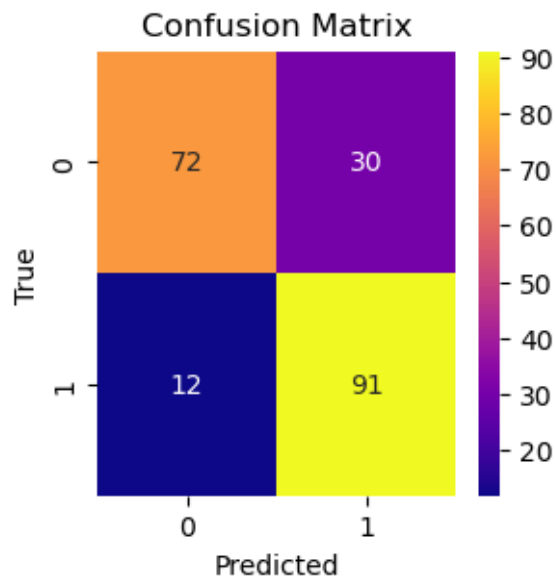
```

print("F1 Score:", f1_score_rf)
print("ROC AUC Score:", roc_auc_rf)

print('Confusion Matrix', conf_matrix_rf)
plt.figure(figsize=(3, 3))
sns.heatmap(conf_matrix, annot=True, cmap='plasma', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

```

Accuracy: 0.975609756097561
 Precision: 0.9803921568627451
 Recall: 0.970873786407767
 Specificity: 0.7058823529411765
 F1 Score: 0.975609756097561
 ROC AUC Score: 0.975632971635256
 Confusion Matrix $\begin{bmatrix} 100 & 2 \\ 3 & 100 \end{bmatrix}$



```

[ ]: #K Nearest Neighbors

from sklearn.neighbors import KNeighborsClassifier

knn_classifier = KNeighborsClassifier()

knn_classifier.fit(X_train, y_train)

```


```
knn_classifier.get_params()
```

```
[ ]: {'algorithm': 'auto',  
      'leaf_size': 30,  
      'metric': 'minkowski',  
      'metric_params': None,  
      'n_jobs': None,  
      'n_neighbors': 5,  
      'p': 2,  
      'weights': 'uniform'}
```

```
[ ]: # K Nearest Neighbors Make predictions on the test set  
y_pred_knn = knn_classifier.predict(X_test)  
  
# Calculate evaluation metrics  
accuracy = accuracy_score(y_test, y_pred_knn)  
precision = precision_score(y_test, y_pred_knn)  
recall = recall_score(y_test, y_pred_knn)  
conf_matrix = confusion_matrix(y_test, y_pred_knn)  
tn, fp, fn, tp = conf_matrix.ravel()  
specificity = tn / (tn + fp)  
sensitivity = tp / (tp + fn)  
npv = tn / (tn + fn)  
f1_score = 2 * (precision * recall) / (precision + recall)  
roc_auc = roc_auc_score(y_test, y_pred_knn)  
  
# Print evaluation metrics  
print("Accuracy:", accuracy)  
print("Precision:", precision)  
print("Recall:", recall)  
print("Specificity:", specificity)  
print("ROC AUC Score:", roc_auc)  
  
print(conf_matrix)  
plt.figure(figsize=(3, 3))  
sns.heatmap(conf_matrix, annot=True, cmap='viridis', fmt='g')  
plt.xlabel('Predicted')  
plt.ylabel('True')  
plt.title('Confusion Matrix')  
plt.show()  
# Plot ROC curve  
# fpr, tpr, thresholds = roc_curve(y_test, y_pred_knn)  
# roc_auc = auc(fpr, tpr)  
  
# plt.figure()
```

```
# plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)'
↳ % roc_auc)
# plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
# plt.xlim([0.0, 1.0])
# plt.ylim([0.0, 1.05])
# plt.xlabel('False Positive Rate')
# plt.ylabel('True Positive Rate')
# plt.title('Receiver Operating Characteristic (ROC) Curve')
# plt.legend(loc="lower right")
# plt.show()
```

Accuracy: 0.8292682926829268
Precision: 0.7931034482758621
Recall: 0.8932038834951457
Specificity: 0.7647058823529411
ROC AUC Score: 0.8289548829240434
[[78 24]
[11 92]]



```
[ ]: #Decision Tree

from sklearn.tree import DecisionTreeClassifier

decision_tree_classifier = DecisionTreeClassifier()

decision_tree_classifier.fit(X_train, y_train)
```

```
[ ]: DecisionTreeClassifier()
```

```
[ ]: # #Decision Tree Make predictions on the test set
y_pred_dt = decision_tree_classifier.predict(X_test)

# Calculate evaluation metrics
accuracy = accuracy_score(y_test, y_pred_dt)
precision = precision_score(y_test, y_pred_dt)
recall = recall_score(y_test, y_pred_dt)
conf_matrix = confusion_matrix(y_test, y_pred_dt)
tn, fp, fn, tp = conf_matrix.ravel()
specificity = tn / (tn + fp)
sensitivity = tp / (tp + fn)
npv = tn / (tn + fn)
f1_score = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_dt)

# Print evaluation metrics
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("Specificity:", specificity)
print("ROC AUC Score:", roc_auc)

print(conf_matrix)
plt.figure(figsize=(3, 3))
sns.heatmap(conf_matrix, annot=True, cmap='inferno', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```

Accuracy: 0.9902439024390244

Precision: 1.0


Recall: 0.9805825242718447

Specificity: 1.0

ROC AUC Score: 0.9902912621359223

```
[[102  0]
```

```
 [ 2 101]]
```



```
[ ]: # Support Vector Machines

from sklearn.svm import SVC

svc_classifier = SVC()

svc_classifier.fit(X_train, y_train)
svc_classifier.get_params()
```

```
[ ]: {'C': 1.0,
      'break_ties': False,
      'cache_size': 200,
      'class_weight': None,
      'coef0': 0.0,
      'decision_function_shape': 'ovr',
      'degree': 3,
      'gamma': 'scale',
      'kernel': 'rbf',
      'max_iter': -1,
      'probability': False,
      'random_state': None,
      'shrinking': True,
      'tol': 0.001,
      'verbose': False}
```

```
[ ]: # Support Vector Machines Make predictions on the test set
y_pred_svc = svc_classifier.predict(X_test)
```

```

# Calculate evaluation metrics
accuracy = accuracy_score(y_test, y_pred_svc)
precision = precision_score(y_test, y_pred_svc)
recall = recall_score(y_test, y_pred_svc)
conf_matrix = confusion_matrix(y_test, y_pred_svc)
tn, fp, fn, tp = conf_matrix.ravel()
specificity = tn / (tn + fp)
f1_score = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_svc)


# Print evaluation metrics
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("Specificity:", specificity)
print("F1 Score:", f1_score)
print("ROC AUC Score:", roc_auc)
print(conf_matrix)
plt.figure(figsize=(3, 3))
sns.heatmap(conf_matrix, annot=True, cmap='magma', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

```

```

Accuracy: 0.8780487804878049
Precision: 0.8482142857142857
Recall: 0.9223300970873787
Specificity: 0.8333333333333334
F1 Score: 0.8837209302325582
ROC AUC Score: 0.8778317152103561
[[85 17]
 [ 8 95]]

```



```
[ ]: lr_false_positive_rate,lr_true_positive_rate,lr_threshold =
    ↳roc_curve(y_test,y_pred_lr)
nb_false_positive_rate,nb_true_positive_rate,nb_threshold =
    ↳roc_curve(y_test,y_pred_nb)
rf_false_positive_rate,rf_true_positive_rate,rf_threshold =
    ↳roc_curve(y_test,y_pred_rf)
knn_false_positive_rate,knn_true_positive_rate,knn_threshold =
    ↳roc_curve(y_test,y_pred_knn)
dt_false_positive_rate,dt_true_positive_rate,dt_threshold =
    ↳roc_curve(y_test,y_pred_dt)
svc_false_positive_rate,svc_true_positive_rate,svc_threshold =
    ↳roc_curve(y_test,y_pred_svc)
sns.set_style('whitegrid')
plt.figure(figsize=(6,4))
plt.title('Receiver Operating Characteristic Curve')

# Plot ROC curves for each classifier with different line styles and colors
plt.plot(lr_false_positive_rate, lr_true_positive_rate, label='Logistic
    ↳Regression', color='blue', linestyle='-')
plt.plot(nb_false_positive_rate, nb_true_positive_rate, label='Naive Bayes',
    ↳color='green', linestyle='-.')
plt.plot(rf_false_positive_rate, rf_true_positive_rate, label='Random Forest',
    ↳color='red', linestyle='--')
plt.plot(knn_false_positive_rate, knn_true_positive_rate, label='K-Nearest
    ↳Neighbor', color='orange', linestyle=':')
```

```


plt.plot(dt_false_positive_rate, dt_true_positive_rate, label='Decision Tree',
        color='purple', linestyle='-')
plt.plot(svc_false_positive_rate, svc_true_positive_rate, label='Support Vector
Classifier', color='cyan', linestyle='-')

# Plot the diagonal reference line
plt.plot([0, 1], [0, 1], ls='--', color='.5')

# Add labels and legend
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.legend()

# Show the plot
plt.show()

```



```

[ ]: # Let's do Hyperparameter tuning!!
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.ensemble import StackingClassifier

[ ]: # Define hyperparameters for logistic regression
param_lr = {
    'C': [0.001, 0.01, 0.1, 1, 10],

```



```

        'solver': ['lbfgs', 'sag', 'saga'],
    }

    # Create grid search using 5-fold cross validation
    grid_search_lr = GridSearchCV(LogisticRegression(max_iter=20000), param_lr,
    ↪cv=5, scoring='accuracy')

    # Perform grid search
    grid_search_lr.fit(X_train, y_train)

    # Print best parameters
    print("Best parameters for Logistic Regression with validation set (Grid_
    ↪Search):", grid_search_lr.best_params_)

```

Best parameters for Logistic Regression with validation set (Grid Search): {'C': 1, 'solver': 'lbfgs'}

```

[ ]: #Logistic Regression: Testing

logistic_regression_best = LogisticRegression(C=0.1, max_iter=10000)
logistic_regression_best.fit(X_train, y_train)

# Before hyperparameter tuning
y_pred_lr_before = logistic_regression.predict(X_test)

accuracy = accuracy_score(y_test, y_pred_lr_before)
precision = precision_score(y_test, y_pred_lr_before)
recall = recall_score(y_test, y_pred_lr_before)
conf_matrix = confusion_matrix(y_test, y_pred_lr_before)
tn, fp, fn, tp = conf_matrix.ravel()
specificity = tn / (tn + fp)
f1_score = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_lr_before)

# Print evaluation metrics
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("Specificity:", specificity)
print("F1 Score:", f1_score)
print("ROC AUC Score:", roc_auc)

logistic_tuned_model = grid_search_lr.best_estimator_
y_pred_lr_best = logistic_tuned_model.predict(X_val)

```

```

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_val, y_pred_lr_best)
precision_after = precision_score(y_val, y_pred_lr_best)
recall_after = recall_score(y_val, y_pred_lr_best)
conf_matrix = confusion_matrix(y_val, y_pred_lr_best)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
f1_score_after = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_val, y_pred_lr_best)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics after Hyperparameter Tuning (Val Set):")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

```

```

Accuracy: 0.7951219512195122
Precision: 0.7520661157024794
Recall: 0.883495145631068
Specificity: 0.7058823529411765
F1 Score: 0.8125000000000001
ROC AUC Score: 0.7946887492861221

```

```

Evaluation Metrics after Hyperparameter Tuning (Val Set):
Accuracy: 0.9073170731707317
Precision: 0.8728813559322034
Recall: 0.9626168224299065
Specificity: 0.8469387755102041
F1 Score: 0.8125000000000001
ROC AUC Score: 0.9047777989700553

```

```

[ ]: # Make predictions on the test set using the tuned logistic regression model
logistic_tuned_model = grid_search_lr.best_estimator_
y_pred_lr_test = logistic_tuned_model.predict(X_test)

# Calculate evaluation metrics for the tuned model on the test set
accuracy_test = accuracy_score(y_test, y_pred_lr_test)
precision_test = precision_score(y_test, y_pred_lr_test)
recall_test = recall_score(y_test, y_pred_lr_test)
conf_matrix_test = confusion_matrix(y_test, y_pred_lr_test)
tn_test, fp_test, fn_test, tp_test = conf_matrix_test.ravel()
specificity_test = tn_test / (tn_test + fp_test)
sensitivity_test = tp_test / (tp_test + fn_test)
npv_test = tn_test / (tn_test + fn_test)

```

```

f1_score_test = 2 * (precision_test * recall_test) / (precision_test +
↪recall_test)
roc_auc_test = roc_auc_score(y_test, y_pred_lr_test)

# Print the evaluation metrics for the tuned model on the test set
print("\nEvaluation Metrics for Tuned Logistic Regression Model on Test Set:")
print("Accuracy:", accuracy_test)
print("Precision:", precision_test)
print("Recall:", recall_test)
print("Specificity:", specificity_test)
print("Sensitivity:", sensitivity_test)
print("NPV:", npv_test)
print("F1 Score:", f1_score_test)
print("ROC AUC Score:", roc_auc_test)

```

Evaluation Metrics for Tuned Logistic Regression Model on Test Set:
Accuracy: 0.7951219512195122
Precision: 0.7520661157024794
Recall: 0.883495145631068
Specificity: 0.7058823529411765
Sensitivity: 0.883495145631068
NPV: 0.8571428571428571
F1 Score: 0.8125000000000001
ROC AUC Score: 0.7946887492861221

```

[ ]: # Random Forest
rf_param_grid = {
    'n_estimators': [50, 100],
    'max_depth': [None, 10],
    'min_samples_split': [2, 5, 7, 10],
    'min_samples_leaf': [1, 2],
    'max_features': ['auto', 'sqrt']
}

grid_search_rf = GridSearchCV(RandomForestClassifier(), rf_param_grid, cv=5,
↪scoring='accuracy')

grid_search_rf.fit(X_train, y_train)

best_params_grid_rf = grid_search_rf.best_params_
best_estimator_grid_rf = grid_search_rf.best_estimator_

print("Results for Random Forest - GridSearchCV:")
print("Best parameters:", best_params_grid_rf)

```

```
print("Best estimator:", best_estimator_grid_rf)
```

D:\Anaconda\envs\workingenv\lib\site-packages\sklearn\model_selection_validation.py:425: FitFailedWarning:
160 fits failed out of a total of 320.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

160 fits failed with the following error:
Traceback (most recent call last):
 File "D:\Anaconda\envs\workingenv\lib\site-packages\sklearn\model_selection_validation.py", line 732, in _fit_and_score
 estimator.fit(X_train, y_train, **fit_params)
 File "D:\Anaconda\envs\workingenv\lib\site-packages\sklearn\base.py", line 1144, in wrapper
 estimator._validate_params()
 File "D:\Anaconda\envs\workingenv\lib\site-packages\sklearn\base.py", line 637, in _validate_params
 validate_parameter_constraints(
 File "D:\Anaconda\envs\workingenv\lib\site-packages\sklearn\utils_param_validation.py", line 95, in
 validate_parameter_constraints
 raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features'
parameter of RandomForestClassifier must be an int in the range [1, inf), a
float in the range (0.0, 1.0], a str among {'log2', 'sqrt'} or None. Got 'auto'
instead.

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
D:\Anaconda\envs\workingenv\lib\site-packages\sklearn\model_selection\_search.py:976: UserWarning: One or more of the
test scores are non-finite: [          nan          nan          nan          nan
nan          nan
          nan          nan          nan          nan          nan          nan
          nan          nan          nan          nan 0.97073171 0.96910569
0.95284553 0.95772358 0.9495935  0.94471545 0.92357724 0.9398374
0.9495935 0.9495935 0.95121951 0.94471545 0.93658537 0.93821138
0.92682927 0.92357724          nan          nan          nan          nan
          nan          nan          nan          nan          nan          nan
          nan          nan          nan          nan          nan          nan
0.95772358 0.97073171 0.95121951 0.95609756 0.94634146 0.94146341
0.93821138 0.94308943 0.93821138 0.95121951 0.94308943 0.94634146
0.93821138 0.93658537 0.92682927 0.9203252 ]
warnings.warn(
```

Results for Random Forest - GridSearchCV:

Best parameters: {'max_depth': None, 'max_features': 'sqrt', 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 50}

Best estimator: RandomForestClassifier(n_estimators=50)

```
[ ]: #Random Forest: Testing

random_forest_best = RandomForestClassifier()
random_forest_best.fit(X_train, y_train)

# Before hyperparameter tuning
y_pred_rf_before = random_forest_best.predict(X_test)

accuracy = accuracy_score(y_test, y_pred_rf_before)
precision = precision_score(y_test, y_pred_rf_before)
recall = recall_score(y_test, y_pred_rf_before)
conf_matrix = confusion_matrix(y_test, y_pred_rf_before)
tn, fp, fn, tp = conf_matrix.ravel()
specificity = tn / (tn + fp)
f1_score = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_rf_before)

# Print evaluation metrics
print("Before HyperParameters Tuning")
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("Specificity:", specificity)
print("F1 Score:", f1_score)
print("ROC AUC Score:", roc_auc)

random_forest_tuned_model = grid_search_rf.best_estimator_
y_pred_rf_best = random_forest_tuned_model.predict(X_val)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_val, y_pred_rf_best)
precision_after = precision_score(y_val, y_pred_rf_best)
recall_after = recall_score(y_val, y_pred_rf_best)
conf_matrix = confusion_matrix(y_val, y_pred_rf_best)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
f1_score_after = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_val, y_pred_rf_best)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics after Hyperparameter Tuning (Val Set):")
```

```

print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

```

Before HyperParameters Tuning

Accuracy: 0.975609756097561

Precision: 0.9803921568627451

Recall: 0.970873786407767

Specificity: 0.9803921568627451

F1 Score: 0.975609756097561

ROC AUC Score: 0.975632971635256

Evaluation Metrics after Hyperparameter Tuning (Val Set):

Accuracy: 0.9951219512195122

Precision: 0.9907407407407407

Recall: 1.0

Specificity: 0.9897959183673469

F1 Score: 0.975609756097561

ROC AUC Score: 0.9948979591836734

```

[ ]: # Make predictions on the test set using the tuned random forest model
random_forest_tuned_model = grid_search_rf.best_estimator_
y_pred_rf_test = random_forest_tuned_model.predict(X_test)

# Calculate evaluation metrics for the tuned model on the test set
accuracy_test = accuracy_score(y_test, y_pred_rf_test)
precision_test = precision_score(y_test, y_pred_rf_test)
recall_test = recall_score(y_test, y_pred_rf_test)
conf_matrix_test = confusion_matrix(y_test, y_pred_rf_test)
tn_test, fp_test, fn_test, tp_test = conf_matrix_test.ravel()
specificity_test = tn_test / (tn_test + fp_test)
sensitivity_test = tp_test / (tp_test + fn_test)
npv_test = tn_test / (tn_test + fn_test)
f1_score_test = 2 * (precision_test * recall_test) / (precision_test +
↪recall_test)
roc_auc_test = roc_auc_score(y_test, y_pred_rf_test)

# Print the evaluation metrics for the tuned model on the test set
print("\nEvaluation Metrics for Tuned Logistic Regression Model on Test Set:")
print("Accuracy:", accuracy_test)
print("Precision:", precision_test)
print("Recall:", recall_test)
print("Specificity:", specificity_test)
print("Sensitivity:", sensitivity_test)

```

```
print("NPV:", npv_test)
print("F1 Score:", f1_score_test)
print("ROC AUC Score:", roc_auc_test)
```

Evaluation Metrics for Tuned Logistic Regression Model on Test Set:

```
Accuracy: 0.975609756097561
Precision: 0.9803921568627451
Recall: 0.970873786407767
Specificity: 0.9803921568627451
Sensitivity: 0.970873786407767
NPV: 0.970873786407767
F1 Score: 0.975609756097561
ROC AUC Score: 0.975632971635256
```

```
[ ]: # K-Nearest Neighbors (KNN)
param_grid_knn = {
    'n_neighbors': range(5, 50),
    'weights': ['uniform', 'distance'],
    'algorithm': ['ball_tree', 'kd_tree', 'brute'],
    'metric' : ['minkowski', 'euclidean', 'manhattan']
}

grid_search_knn = GridSearchCV(KNeighborsClassifier(), param_grid_knn, cv=5,
    ↪scoring='accuracy')

grid_search_knn.fit(X_train, y_train)

best_params_grid_knn = grid_search_knn.best_params_
best_estimator_grid_knn = grid_search_knn.best_estimator_

print("Results for K Nearest Neighbors - GridSearchCV:")
print("Best parameters:", best_params_grid_knn)
print("Best estimator:", best_estimator_grid_knn)
```

Results for K Nearest Neighbors - GridSearchCV:

```
Best parameters: {'algorithm': 'ball_tree', 'metric': 'manhattan',
'n_neighbors': 44, 'weights': 'distance'}
Best estimator: KNeighborsClassifier(algorithm='ball_tree', metric='manhattan',
n_neighbors=44,
                                weights='distance')
```

```
[ ]: #K Nearest Neighbors: Testing

# Before hyperparameter tuning
y_pred_knn_before = knn_classifier.predict(X_test)

accuracy = accuracy_score(y_test, y_pred_knn_before)
```

```

precision = precision_score(y_test, y_pred_knn_before)
recall = recall_score(y_test, y_pred_knn_before)
conf_matrix = confusion_matrix(y_test, y_pred_knn_before)
tn, fp, fn, tp = conf_matrix.ravel()
specificity = tn / (tn + fp)
f1_score = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_knn_before)

# Print evaluation metrics
print("Before Hyper Parameters Tuning")
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("Specificity:", specificity)
print("F1 Score:", f1_score)
print("ROC AUC Score:", roc_auc)

# After hyperparameter tuning
# Make predictions on the test set using tuned model

knn_tuned_model = grid_search_knn.best_estimator_
y_pred_knn_after_gridSearch = knn_tuned_model.predict(X_val)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_val, y_pred_knn_after_gridSearch)
precision_after = precision_score(y_val, y_pred_knn_after_gridSearch)
recall_after = recall_score(y_val, y_pred_knn_after_gridSearch)
conf_matrix = confusion_matrix(y_val, y_pred_knn_after_gridSearch)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
f1_score_after = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_val, y_pred_knn_after_gridSearch)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics after Hyperparameter Tuning (GridSearch):")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

```

Before Hyper Parameters Tuning
 Accuracy: 0.8292682926829268
 Precision: 0.7931034482758621
 Recall: 0.8932038834951457
 Specificity: 0.7647058823529411
 F1 Score: 0.8401826484018265

ROC AUC Score: 0.8289548829240434

Evaluation Metrics after Hyperparameter Tuning (GridSearch):

Accuracy: 0.9804878048780488

Precision: 0.9904761904761905

Recall: 0.9719626168224299

Specificity: 0.9897959183673469

F1 Score: 0.8401826484018265

ROC AUC Score: 0.9808792675948884

```
[ ]: # After hyperparameter tuning
      # Make predictions on the test set using tuned model

knn_tuned_model_test = grid_search_knn.best_estimator_
y_pred_knn_after_gridSearch_test = knn_tuned_model_test.predict(X_test)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_test, y_pred_knn_after_gridSearch_test)
precision_after = precision_score(y_test, y_pred_knn_after_gridSearch_test)
recall_after = recall_score(y_test, y_pred_knn_after_gridSearch_test)
conf_matrix = confusion_matrix(y_test, y_pred_knn_after_gridSearch_test)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
sensitivity_after = tp / (tp + fn)
npv_after = tn / (tn + fn)
f1_score_after = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_knn_after_gridSearch_test)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics after Hyperparameter Tuning (GridSearch):")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("Sensitivity:", sensitivity_after)
print("NPV:", npv_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)
```

Evaluation Metrics after Hyperparameter Tuning (GridSearch):

Accuracy: 0.9902439024390244

Precision: 0.9809523809523809

Recall: 1.0

Specificity: 0.9803921568627451

Sensitivity: 1.0

NPV: 1.0

F1 Score: 0.8401826484018265

ROC AUC Score: 0.9901960784313725

```
[ ]: # Decision Trees:
      # Define the parameter grid

param_grid = {
    'max_depth': [2, 10, 20, 30],
    'min_samples_split': [2, 4, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt', 'log2']
}
# Initialize the decision tree classifier
decision_tree = DecisionTreeClassifier(random_state=42)

# Initialize GridSearchCV
grid_search_decision_tree = GridSearchCV(estimator=decision_tree,
    param_grid=param_grid, cv=5, scoring='accuracy')

# Perform hyperparameter tuning
grid_search_decision_tree.fit(X_train, y_train)

# Best parameters found during grid search
best_params = grid_search_decision_tree.best_params_
print("Best Parameters:", best_params)
```

Best Parameters: {'max_depth': 20, 'max_features': 'sqrt', 'min_samples_leaf': 1, 'min_samples_split': 2}

D:\Anaconda\envs\workingenv\lib\site-packages\sklearn\model_selection_validation.py:425: FitFailedWarning:
180 fits failed out of a total of 540.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

```
-----
180 fits failed with the following error:
Traceback (most recent call last):
  File "D:\Anaconda\envs\workingenv\lib\site-packages\sklearn\model_selection\_validation.py", line 732, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "D:\Anaconda\envs\workingenv\lib\site-packages\sklearn\base.py", line 1144, in wrapper
    estimator._validate_params()
  File "D:\Anaconda\envs\workingenv\lib\site-packages\sklearn\base.py", line 637, in _validate_params
```

```

    validate_parameter_constraints(
    File "D:\Anaconda\envs\workingenv\lib\site-
packages\sklearn\utils\_param_validation.py", line 95, in
validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features'
parameter of DecisionTreeClassifier must be an int in the range [1, inf), a
float in the range (0.0, 1.0], a str among {'log2', 'sqrt'} or None. Got 'auto'
instead.

warnings.warn(some_fits_failed_message, FitFailedWarning)
D:\Anaconda\envs\workingenv\lib\site-
packages\sklearn\model_selection\_search.py:976: UserWarning: One or more of the
test scores are non-finite: [      nan      nan      nan      nan
nan      nan
      nan      nan      nan 0.68292683 0.68292683 0.68292683
0.68292683 0.68292683 0.68292683 0.68292683 0.68292683 0.68292683
0.68292683 0.68292683 0.68292683      nan      nan      nan
      nan      nan      nan      nan      nan      nan
0.93821138 0.91869919 0.86829268 0.89593496 0.89593496 0.85528455
0.89105691 0.89105691 0.8601626 0.93821138 0.91869919 0.86829268
0.89593496 0.89593496 0.85528455 0.89105691 0.89105691 0.8601626
      nan      nan      nan      nan      nan      nan
      nan      nan      nan 0.95121951 0.91869919 0.86829268
0.90731707 0.90731707 0.85528455 0.89105691 0.89105691 0.8601626
0.95121951 0.91869919 0.86829268 0.90731707 0.90731707 0.85528455
0.89105691 0.89105691 0.8601626      nan      nan      nan
      nan      nan      nan      nan      nan      nan
0.95121951 0.91869919 0.86829268 0.90731707 0.90731707 0.85528455
0.89105691 0.89105691 0.8601626 0.95121951 0.91869919 0.86829268
0.90731707 0.90731707 0.85528455 0.89105691 0.89105691 0.8601626 ]
warnings.warn(

```

```
[ ]: #Decision Trees: Testing
```

```

# Before hyperparameter tuning
y_pred_dt_before = decision_tree_classifier.predict(X_test)

accuracy = accuracy_score(y_test, y_pred_dt_before)
precision = precision_score(y_test, y_pred_dt_before)
recall = recall_score(y_test, y_pred_dt_before)
conf_matrix = confusion_matrix(y_test, y_pred_dt_before)
tn, fp, fn, tp = conf_matrix.ravel()
specificity = tn / (tn + fp)
sensitivity = tp / (tp + fn)
npv = tn / (tn + fn)

```

```

f1_score = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_dt_before)

# Print evaluation metrics
print("Before Hyper parameters tuning")
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("Specificity:", specificity)
print("F1 Score:", f1_score)
print("ROC AUC Score:", roc_auc)

# After hyperparameter tuning
# Make predictions on the test set using tuned model

decision_tree_tuned_model = grid_search_decision_tree.best_estimator_
y_pred_dt_after_gridSearch = decision_tree_tuned_model.predict(X_val)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_val, y_pred_dt_after_gridSearch)
precision_after = precision_score(y_val, y_pred_dt_after_gridSearch)
recall_after = recall_score(y_val, y_pred_dt_after_gridSearch)
conf_matrix = confusion_matrix(y_val, y_pred_dt_after_gridSearch)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
f1_score_after = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_val, y_pred_dt_after_gridSearch)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics after Hyperparameter Tuning (GridSearch):")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

```

Before Hyper parameters tuning
 Accuracy: 0.9707317073170731
 Precision: 0.9619047619047619
 Recall: 0.9805825242718447
 Specificity: 0.9607843137254902
 F1 Score: 0.9711538461538461
 ROC AUC Score: 0.9706834189986675

Evaluation Metrics after Hyperparameter Tuning (GridSearch):
 Accuracy: 0.9902439024390244
 Precision: 0.981651376146789

Recall: 1.0
Specificity: 0.9795918367346939
F1 Score: 0.9711538461538461
ROC AUC Score: 0.9897959183673469

```
[ ]: # After hyperparameter tuning
# Make predictions on the test set using tuned model

decision_tree_tuned_model = grid_search_decision_tree.best_estimator_
y_pred_dt_after_gridSearch_test = decision_tree_tuned_model.predict(X_test)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_test, y_pred_dt_after_gridSearch_test)
precision_after = precision_score(y_test, y_pred_dt_after_gridSearch_test)
recall_after = recall_score(y_test, y_pred_dt_after_gridSearch_test)
conf_matrix = confusion_matrix(y_test, y_pred_dt_after_gridSearch_test)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
sensitivity_after = tp / (tp + fn)
npv_after = tn / (tn + fn)
f1_score_after = 2 * (precision_after * recall_after) / (precision_after +
↪recall_after)
roc_auc = roc_auc_score(y_test, y_pred_dt_after_gridSearch_test)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics after Hyperparameter Tuning (GridSearch):")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("Sensitivity:", sensitivity_after)
print("NPV:", npv_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)
```

Evaluation Metrics after Hyperparameter Tuning (GridSearch):
Accuracy: 0.9609756097560975
Precision: 0.9523809523809523
Recall: 0.970873786407767
Specificity: 0.9509803921568627
Sensitivity: 0.970873786407767
NPV: 0.97
F1 Score: 0.9615384615384616
ROC AUC Score: 0.9609270892823149

```
[ ]: # Support Vector Machine (SVM)
param_grid_svc = {
```

```

        'C': [0.001, 0.01, 0.1, 1, 10, 100],
        'kernel': ['linear', 'poly', 'rbf', 'sigmoid'],
        'gamma': ['scale', 'auto']
    }

    grid_search_svc = GridSearchCV(SVC(), param_grid_svc, cv=5, scoring='accuracy')
    # random_search_svc = RandomizedSearchCV(SVC(), param_grid_svc, cv=5,
    # ↪n_iter=100, scoring='accuracy')

    grid_search_svc.fit(X_train, y_train)
    # random_search_svc.fit(X_val, y_val)

    best_params_grid_svc = grid_search_svc.best_params_
    best_estimator_grid_svc = grid_search_svc.best_estimator_

    # best_params_random_svc = random_search_svc.best_params_
    # best_estimator_random_svc = random_search_svc.best_estimator_

    print("Results for Support Vector Machines - GridSearchCV:")
    print("Best parameters:", best_params_grid_svc)
    print("Best estimator:", best_estimator_grid_svc)

    # print("\nResults for Support Vector Machines - RandomizedSearchCV:")
    # print("Best parameters:", best_params_random_svc)
    # print("Best estimator:", best_estimator_random_svc)

```

Results for Support Vector Machines - GridSearchCV:
 Best parameters: {'C': 10, 'gamma': 'scale', 'kernel': 'rbf'}
 Best estimator: SVC(C=10)

```

[ ]: #Support Vector Machines: Testing

# Before hyperparameter tuning
y_pred_svc_before = svc_classifier.predict(X_test)

accuracy = accuracy_score(y_test, y_pred_svc_before)
precision = precision_score(y_test, y_pred_svc_before)
recall = recall_score(y_test, y_pred_svc_before)
conf_matrix = confusion_matrix(y_test, y_pred_svc_before)
tn, fp, fn, tp = conf_matrix.ravel()
specificity = tn / (tn + fp)
f1_score = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_svc_before)

# Print evaluation metrics
print('Before Hyper Parameters tuning')
print("Accuracy:", accuracy)

```

```

print("Precision:", precision)
print("Recall:", recall)
print("Specificity:", specificity)
print("F1 Score:", f1_score)
print("ROC AUC Score:", roc_auc)

# After hyperparameter tuning
# Make predictions on the test set using tuned mode
best_svm_grid = grid_search_svc.best_estimator_
y_test_svc_after_gridSearch = best_svm_grid.predict(X_val)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_val, y_test_svc_after_gridSearch)
precision_after = precision_score(y_val, y_test_svc_after_gridSearch)
recall_after = recall_score(y_val, y_test_svc_after_gridSearch)
conf_matrix = confusion_matrix(y_val, y_test_svc_after_gridSearch)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
f1_score_after = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_val, y_test_svc_after_gridSearch)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics after Hyperparameter Tuning (GridSearch):")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

```

Before Hyper Parameters tuning
Accuracy: 0.8780487804878049
Precision: 0.8482142857142857
Recall: 0.9223300970873787
Specificity: 0.8333333333333334
F1 Score: 0.8837209302325582
ROC AUC Score: 0.8778317152103561

Evaluation Metrics after Hyperparameter Tuning (GridSearch):
Accuracy: 0.975609756097561
Precision: 0.9811320754716981
Recall: 0.9719626168224299
Specificity: 0.9795918367346939
F1 Score: 0.8837209302325582
ROC AUC Score: 0.9757772267785619

```
[ ]: # After hyperparameter tuning
# Make predictions on the test set using tuned mode
best_svm_grid_test = grid_search_svc.best_estimator_
y_svc_after_gridSearch_test = best_svm_grid_test.predict(X_test)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_test, y_svc_after_gridSearch_test)
precision_after = precision_score(y_test, y_svc_after_gridSearch_test)
recall_after = recall_score(y_test, y_svc_after_gridSearch_test)
conf_matrix = confusion_matrix(y_test, y_svc_after_gridSearch_test)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
sensitivity_after = tp / (tp + fn)
npv_after = tn / (tn + fn)
f1_score_after = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_svc_after_gridSearch_test)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics after Hyperparameter Tuning (GridSearch):")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("Sensitivity:", sensitivity_after)
print("NPV:", npv_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)
```

Evaluation Metrics after Hyperparameter Tuning (GridSearch):

Accuracy: 0.9512195121951219
Precision: 0.9603960396039604
Recall: 0.941747572815534
Specificity: 0.9607843137254902
Sensitivity: 0.941747572815534
NPV: 0.9423076923076923
F1 Score: 0.8837209302325582
ROC AUC Score: 0.9512659432705121

```
[ ]: lr_false_positive_rate,lr_true_positive_rate,lr_threshold =_
    ↪roc_curve(y_test,y_pred_lr_test)
nb_false_positive_rate,nb_true_positive_rate,nb_threshold =_
    ↪roc_curve(y_test,y_pred_nb)
rf_false_positive_rate,rf_true_positive_rate,rf_threshold =_
    ↪roc_curve(y_test,y_pred_rf_test)
knn_false_positive_rate,knn_true_positive_rate,knn_threshold =_
    ↪roc_curve(y_test,y_pred_knn_after_gridSearch_test)
```



```

dt_false_positive_rate,dt_true_positive_rate,dt_threshold =
    ↪roc_curve(y_test,y_pred_dt_after_gridSearch_test)
svc_false_positive_rate,svc_true_positive_rate,svc_threshold =
    ↪roc_curve(y_test,y_svc_after_gridSearch_test)
sns.set_style('whitegrid')
plt.figure(figsize=(10,8))
plt.title('Receiver Operating Characteristic Curve')


# Plot ROC curves for each classifier with different line styles and colors
plt.plot(lr_false_positive_rate, lr_true_positive_rate, label='Logistic
    ↪Regression', color='blue', linestyle='-')
plt.plot(nb_false_positive_rate, nb_true_positive_rate, label='Naive Bayes',
    ↪color='green', linestyle='-.')
plt.plot(rf_false_positive_rate, rf_true_positive_rate, label='Random Forest',
    ↪color='red', linestyle='--')
plt.plot(knn_false_positive_rate, knn_true_positive_rate, label='K-Nearest
    ↪Neighbor', color='orange', linestyle=':')
plt.plot(dt_false_positive_rate, dt_true_positive_rate, label='Decision Tree',
    ↪color='purple', linestyle='-')
plt.plot(svc_false_positive_rate, svc_true_positive_rate, label='Support Vector
    ↪Classifier', color='cyan', linestyle='-')

# Plot the diagonal reference line
plt.plot([0, 1], [0, 1], ls='--', color='.5')

# Add labels and legend
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.legend()

# Show the plot
plt.show()

```



```
[ ]: # LETS DO Feature Selection
from sklearn.feature_selection import RFECV
from sklearn.model_selection import StratifiedKFold

[ ]: #Logistic Regression
# Wrapper-based feature selection with Recursive Feature Elimination (RFE) for
↳ Logistic Regression

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

logistic_regression_rfecv = RFECV(estimator=logistic_regression_best, step=1,
↳ cv=cv, scoring='accuracy')

logistic_regression_rfecv.fit(X_train, y_train)

print("Optimal number of features:", logistic_regression_rfecv.n_features_)
```

```

X_train_selected_wrapper = logistic_regression_rfecv.fit_transform(X_train,
↪y_train)

selected_features = X.columns[logistic_regression_rfecv.get_support()]

print("Selected features for Random Forest (Wrapper-Based):", selected_features)

```

Optimal number of features: 6

Selected features for Random Forest (Wrapper-Based): Index(['sex', 'cp', 'thalach', 'oldpeak', 'ca', 'thal'], dtype='object')

```

[ ]: # After hyperparameter tuning
y_pred_lr_best = logistic_regression_best.predict(X_test)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_test, y_pred_lr_best)
precision_after = precision_score(y_test, y_pred_lr_best)
recall_after = recall_score(y_test, y_pred_lr_best)
conf_matrix = confusion_matrix(y_test, y_pred_lr_best)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
f1_score_after = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_lr_best)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics before Feature Selection:")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

# After feature selection
# X_test_selected_features = logistic_regression_rfecv.transform(X_test)

y_pred_lr_filtered = logistic_regression_rfecv.predict(X_test)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_test, y_pred_lr_filtered)
precision_after = precision_score(y_test, y_pred_lr_filtered)
recall_after = recall_score(y_test, y_pred_lr_filtered)
conf_matrix_test = confusion_matrix(y_test, y_pred_lr_filtered)
tn, fp, fn, tp = conf_matrix_test.ravel()
specificity_after = tn / (tn + fp)
f1_score_after = 2 * (precision * recall) / (precision + recall)

```

```

roc_auc = roc_auc_score(y_test, y_pred_lr_filtered)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics after Feature Selection:")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

print(conf_matrix_test)

plt.figure(figsize=(3, 3))
sns.heatmap(conf_matrix_test, annot=True, cmap='Purples', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()


```

Evaluation Metrics before Feature Selection:

Accuracy: 0.7951219512195122
Precision: 0.7520661157024794
Recall: 0.883495145631068
Specificity: 0.7058823529411765
F1 Score: 0.8837209302325582
ROC AUC Score: 0.7946887492861221

Evaluation Metrics after Feature Selection:

Accuracy: 0.8097560975609757
Precision: 0.7758620689655172
Recall: 0.8737864077669902
Specificity: 0.7058823529411765
F1 Score: 0.8837209302325582
ROC AUC Score: 0.8094422234913383
[[76 26]
[13 90]]



```
[ ]: # Naive Bayes

from sklearn.feature_selection import SelectKBest, mutual_info_classif
from sklearn.preprocessing import MinMaxScaler

# Apply Min-Max scaling to ensure non-negative values
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)

naive_bayes_selector = SelectKBest(score_func=mutual_info_classif, k=8)

X_train_selected = naive_bayes_selector.fit_transform(X_train_scaled, y_train)

nb_filtered = GaussianNB()
nb_filtered.fit(X_train_selected, y_train)

selected_features_mask = naive_bayes_selector.get_support()

# Count the number of selected features
num_selected_features = np.sum(selected_features_mask)

print("Number of selected features:", num_selected_features)

selected_feature_indices = naive_bayes_selector.get_support(indices=True)
```

```
# Get the names of selected features
selected_features = X.columns[selected_feature_indices]

print("Selected features:", selected_features)
```

Number of selected features: 8
Selected features: Index(['cp', 'trestbps', 'chol', 'thalach', 'exang', 'oldpeak', 'ca', 'thal'], dtype='object')

```
[ ]: # Naive Bayes Model
y_pred_nb = naive_bayes_classifier.predict(X_test)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_test, y_pred_nb)
precision_after = precision_score(y_test, y_pred_nb)
recall_after = recall_score(y_test, y_pred_nb)
conf_matrix = confusion_matrix(y_test, y_pred_nb)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
sensitivity_after = tp / (tp + fn)
npv_after = tn / (tn + fn)
f1_score_after = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_nb)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics before Feature Selection:")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

# After feature selection
X_test_scaled = scaler.fit_transform(X_test)
X_test_filtered = naive_bayes_selector.transform(X_test_scaled)
y_pred_nb_filtered = nb_filtered.predict(X_test_filtered)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_test, y_pred_nb_filtered)
precision_after = precision_score(y_test, y_pred_nb_filtered)
recall_after = recall_score(y_test, y_pred_nb_filtered)
conf_matrix_test = confusion_matrix(y_test, y_pred_nb_filtered)
tn, fp, fn, tp = conf_matrix_test.ravel()
specificity_after = tn / (tn + fp)
sensitivity_after = tp / (tp + fn)
```

```

npv_after = tn / (tn + fn)
f1_score_after = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_nb_filtered)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics after Feature Selection:")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

print(conf_matrix_test)

plt.figure(figsize=(3, 3))
sns.heatmap(conf_matrix_test, annot=True, cmap='Greens', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()


```

Evaluation Metrics before Feature Selection:

Accuracy: 0.8048780487804879
Precision: 0.7560975609756098
Recall: 0.9029126213592233
Specificity: 0.7058823529411765
F1 Score: 0.8837209302325582
ROC AUC Score: 0.8043974871501998

Evaluation Metrics after Feature Selection:

Accuracy: 0.7951219512195122
Precision: 0.7563025210084033
Recall: 0.8737864077669902
Specificity: 0.7058823529411765
F1 Score: 0.8837209302325582
ROC AUC Score: 0.7947363411383972
[[73 29]
[13 90]]



```
[ ]: # K Nearest Neighbors

from sklearn.feature_selection import SelectKBest, mutual_info_classif

knn_selector = SelectKBest(score_func=mutual_info_classif, k=10)

X_train_selected = knn_selector.fit_transform(X_train, y_train)

# Train KNeighborsClassifier using the selected features

grid_search_knn_filtered = copy.copy(grid_search_knn)
best_params_knn = grid_search_knn_filtered.best_params_

knn_classifier_filtered = KNeighborsClassifier(**best_params_knn)
knn_classifier_filtered.fit(X_train_selected, y_train)

selected_features_mask = knn_selector.get_support()

# Count the number of selected features
num_selected_features = np.sum(selected_features_mask)

print("Number of selected features:", num_selected_features)

selected_feature_indices = knn_selector.get_support(indices=True)
```



```

# Get the names of selected features
selected_features = X.columns[selected_feature_indices]

print("Selected features:", selected_features)

```

Number of selected features: 10

```

Selected features: Index(['age', 'cp', 'trestbps', 'chol', 'thalach', 'exang',
                        'oldpeak', 'slope',
                        'ca', 'thal'],
                        dtype='object')

```

```

[ ]: # After hyperparameter tuning
y_pred_knn_after_gridSearch = knn_tuned_model.predict(X_test)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_test, y_pred_knn_after_gridSearch)
precision_after = precision_score(y_test, y_pred_knn_after_gridSearch)
recall_after = recall_score(y_test, y_pred_knn_after_gridSearch)
conf_matrix = confusion_matrix(y_test, y_pred_knn_after_gridSearch)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
sensitivity_after = tp / (tp + fn)
npv_after = tn / (tn + fn)
f1_score_after = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_knn_after_gridSearch)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics after Hyperparameter Tuning:")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

# After feature selection
X_test_selected = knn_selector.transform(X_test)
print(X_test_selected.shape)
y_pred_knn_filtered = knn_classifier_filtered.predict(X_test_selected)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_test, y_pred_knn_filtered)
precision_after = precision_score(y_test, y_pred_knn_filtered)
recall_after = recall_score(y_test, y_pred_knn_filtered)
conf_matrix_test = confusion_matrix(y_test, y_pred_knn_filtered)
tn, fp, fn, tp = conf_matrix_test.ravel()
specificity_after = tn / (tn + fp)

```

```

sensitivity_after = tp / (tp + fn)
npv_after = tn / (tn + fn)
f1_score_after = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_knn_filtered)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics before Feature Selection:")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

print(conf_matrix_test)

plt.figure(figsize=(3, 3))
sns.heatmap(conf_matrix_test, annot=True, cmap='YlOrBr', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()


```

Evaluation Metrics after Hyperparameter Tuning:

Accuracy: 0.9902439024390244
Precision: 0.9809523809523809
Recall: 1.0
Specificity: 0.9803921568627451
F1 Score: 0.8837209302325582
ROC AUC Score: 0.9901960784313725
(205, 10)

Evaluation Metrics before Feature Selection:

Accuracy: 0.9902439024390244
Precision: 0.9809523809523809
Recall: 1.0
Specificity: 0.9803921568627451
F1 Score: 0.8837209302325582
ROC AUC Score: 0.9901960784313725
[[100 2]
[0 103]]



```
[ ]: # Random Forest: Feature Selection
cv = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)

random_forest_filtering = RandomForestClassifier()

random_forest_rfecv = RFECV(estimator=random_forest_classifier, step=1, cv=cv,
    ↳scoring='accuracy')

random_forest_rfecv.fit(X_train, y_train)

print("Optimal number of features:", random_forest_rfecv.n_features_)

X_train_selected_wrapper = random_forest_rfecv.fit_transform(X_train, y_train)

selected_features = X.columns[random_forest_rfecv.get_support()]

print("Selected features for Random Forest (Wrapper-Based):", selected_features)
```

```
Optimal number of features: 8
Selected features for Random Forest (Wrapper-Based): Index(['age', 'cp',
'trestbps', 'chol', 'thalach', 'oldpeak', 'ca', 'thal'], dtype='object')
```

```
[ ]: # After hyperparameter tuning
y_pred_random_forest_best = random_forest_classifier.predict(X_test)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_test, y_pred_random_forest_best)
```

```

precision_after = precision_score(y_test, y_pred_random_forest_best)
recall_after = recall_score(y_test, y_pred_random_forest_best)
conf_matrix = confusion_matrix(y_test, y_pred_random_forest_best)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
sensitivity_after = tp / (tp + fn)
npv_after = tn / (tn + fn)
f1_score_after = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_random_forest_best)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics before Feature Selection:")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

# After feature selection
# X_test_selected_features = logistic_regression_rfecv.transform(X_test)

y_pred_random_forest_filtered = random_forest_rfecv.predict(X_test)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_test, y_pred_random_forest_filtered)
precision_after = precision_score(y_test, y_pred_random_forest_filtered)
recall_after = recall_score(y_test, y_pred_random_forest_filtered)
conf_matrix_test = confusion_matrix(y_test, y_pred_random_forest_filtered)
tn, fp, fn, tp = conf_matrix_test.ravel()
specificity_after = tn / (tn + fp)
sensitivity_after = tp / (tp + fn)
npv_after = tn / (tn + fn)
f1_score_after = 2 * (precision_after * recall_after) / (precision_after +
↵recall_after)
roc_auc = roc_auc_score(y_test, y_pred_random_forest_filtered)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics after Feature Selection:")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

```

```

print(conf_matrix_test)

plt.figure(figsize=(3, 3))
sns.heatmap(conf_matrix_test, annot=True, cmap='PuRd', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

```

Evaluation Metrics before Feature Selection:

Accuracy: 0.975609756097561

Precision: 0.9803921568627451

Recall: 0.970873786407767

Specificity: 0.9803921568627451

F1 Score: 0.8837209302325582

ROC AUC Score: 0.975632971635256

Evaluation Metrics after Feature Selection:

Accuracy: 0.975609756097561

Precision: 0.9803921568627451

Recall: 0.970873786407767


Specificity: 0.9803921568627451

F1 Score: 0.975609756097561

ROC AUC Score: 0.975632971635256

[[100 2]

[3 100]]



```
[ ]: # Decision Trees: Feature Selection

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

decision_trees_rfecv = RFECV(estimator=decision_tree_classifier, step=1, cv=cv,
    ↳scoring='accuracy')

decision_trees_rfecv.fit(X_train, y_train)

print("Optimal number of features:", decision_trees_rfecv.n_features_)

X_train_selected_wrapper = decision_trees_rfecv.fit_transform(X_train, y_train)

selected_features = X.columns[decision_trees_rfecv.get_support()]

print("Selected features for Decision Trees (Wrapper-Based):",
    ↳selected_features)
```

```
Optimal number of features: 8
Selected features for Decision Trees (Wrapper-Based): Index(['age', 'sex', 'cp',
'trestbps', 'chol', 'fbs', 'restecg', 'thalach',
'exang', 'oldpeak', 'slope', 'ca', 'thal'],
dtype='object')
```

```
[ ]: # After hyperparameter tuning
y_pred_decision_best = decision_tree_classifier.predict(X_test)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_test, y_pred_decision_best)
precision_after = precision_score(y_test, y_pred_decision_best)
recall_after = recall_score(y_test, y_pred_decision_best)
conf_matrix = confusion_matrix(y_test, y_pred_decision_best)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
sensitivity_after = tp / (tp + fn)
npv_after = tn / (tn + fn)
f1_score_after = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_decision_best)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics before Feature Selection:")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)
```

```

# After feature selection
# X_test_selected_features = logistic_regression_rfecv.transform(X_test)

y_pred_decision_trees_filtered = decision_trees_rfecv.predict(X_test)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_test, y_pred_decision_trees_filtered)
precision_after = precision_score(y_test, y_pred_decision_trees_filtered)
recall_after = recall_score(y_test, y_pred_decision_trees_filtered)
conf_matrix_test = confusion_matrix(y_test, y_pred_decision_trees_filtered)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
sensitivity_after = tp / (tp + fn)
npv_after = tn / (tn + fn)
f1_score_after = 2 * (precision_after * recall_after) / (precision_after +
↪recall_after)
roc_auc = roc_auc_score(y_test, y_pred_decision_trees_filtered)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics after Feature Selection:")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

print(conf_matrix_test)

plt.figure(figsize=(3, 3))
sns.heatmap(conf_matrix_test, annot=True, cmap='RdPu', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

```

Evaluation Metrics before Feature Selection:

Accuracy: 0.9707317073170731
Precision: 0.9619047619047619
Recall: 0.9805825242718447
Specificity: 0.9607843137254902
F1 Score: 0.8837209302325582
ROC AUC Score: 0.9706834189986675

Evaluation Metrics after Feature Selection:

Accuracy: 0.9804878048780488

Precision: 0.9805825242718447

Recall: 0.9805825242718447


Specificity: 0.9607843137254902

F1 Score: 0.9805825242718447

ROC AUC Score: 0.9804873405672949

```
[[100  2]
```

```
 [ 2 101]]
```



```
[ ]: # Support Vector Machines: Feature Selection
svm_selector = SelectKBest(score_func=mutual_info_classif, k=13)
X_train_selected = svm_selector.fit_transform(X_train, y_train)

feature_scores = svm_selector.scores_

# Get the p-values of all features (if applicable)
# p_values = svm_selector.pvalues_ # Uncomment if using a scoring function
# that supports p-values

# Get the indices of the selected features
selected_feature_indices = svm_selector.get_support(indices=True)

# Get the names of the selected features
selected_features = X.columns[selected_feature_indices]

# Print details
```



```

print("Selected Features:")
for feature, score in zip(selected_features,
    ↪ feature_scores[selected_feature_indices]):
    print(f"Feature: {feature}, Score: {score}")
# Train SVM Classifier using the selected features

grid_search_svm_filtered = copy.copy(grid_search_svc)
best_params_svm = grid_search_svm_filtered.best_params_

svm_classifier_filtered = SVC(**best_params_svm)
svm_classifier_filtered.fit(X_train_selected, y_train)

```

Selected Features:

```

Feature: age, Score: 0.11943488155505566
Feature: sex, Score: 0.04326817650116621
Feature: cp, Score: 0.14450529811040713
Feature: trestbps, Score: 0.0623124022070245
Feature: chol, Score: 0.1785400598151876
Feature: fbs, Score: 0.0
Feature: restecg, Score: 0.02645413897071247
Feature: thalach, Score: 0.15982183438875674
Feature: exang, Score: 0.08456900931628786
Feature: oldpeak, Score: 0.16354731339701334
Feature: slope, Score: 0.08910389414085573
Feature: ca, Score: 0.1147388091937962
Feature: thal, Score: 0.12563386431926427

```

[]: SVC(C=10)

```

[ ]: # After hyperparameter tuning
y_pred_svm_best = best_svm_grid_test.predict(X_test)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_test, y_pred_svm_best)
precision_after = precision_score(y_test, y_pred_svm_best)
recall_after = recall_score(y_test, y_pred_svm_best)
conf_matrix = confusion_matrix(y_test, y_pred_svm_best)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
sensitivity_after = tp / (tp + fn)
npv_after = tn / (tn + fn)
f1_score_after = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_svm_best)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics before Feature Selection:")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)

```

```

print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

# After feature selection
# X_test_selected_features = logistic_regression_rfecv.transform(X_test)

X_test_filtered = svm_selector.transform(X_test)
y_pred_svm_filtered = svm_classifier_filtered.predict(X_test_filtered)

# Calculate evaluation metrics after tuning
accuracy_after = accuracy_score(y_test, y_pred_svm_filtered)
precision_after = precision_score(y_test, y_pred_svm_filtered)
recall_after = recall_score(y_test, y_pred_svm_filtered)
conf_matrix_test = confusion_matrix(y_test, y_pred_svm_filtered)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
sensitivity_after = tp / (tp + fn)
npv_after = tn / (tn + fn)
f1_score_after = 2 * (precision_after * recall_after) / (precision_after +
    ↪ recall_after)
roc_auc = roc_auc_score(y_test, y_pred_svm_filtered)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics after Feature Selection:")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)

print(conf_matrix_test)

plt.figure(figsize=(3, 3))
sns.heatmap(conf_matrix_test, annot=True, cmap='BuPu', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

```

Evaluation Metrics before Feature Selection:

Accuracy: 0.9512195121951219

Precision: 0.9603960396039604


Recall: 0.941747572815534

Specificity: 0.9607843137254902
F1 Score: 0.8837209302325582
ROC AUC Score: 0.9512659432705121

Evaluation Metrics after Feature Selection:

Accuracy: 0.9512195121951219
Precision: 0.9603960396039604
Recall: 0.941747572815534
Specificity: 0.9607843137254902
F1 Score: 0.9509803921568628
ROC AUC Score: 0.9512659432705121

```
[[98  4]  
 [ 6 97]]
```



```
[ ]: lr_false_positive_rate,lr_true_positive_rate,lr_threshold =_
      ↪roc_curve(y_test,y_pred_lr_filtered)
nb_false_positive_rate,nb_true_positive_rate,nb_threshold =_
      ↪roc_curve(y_test,y_pred_nb_filtered)
rf_false_positive_rate,rf_true_positive_rate,rf_threshold =_
      ↪roc_curve(y_test,y_pred_random_forest_filtered)
knn_false_positive_rate,knn_true_positive_rate,knn_threshold =_
      ↪roc_curve(y_test,y_pred_knn_filtered)
dt_false_positive_rate,dt_true_positive_rate,dt_threshold =_
      ↪roc_curve(y_test,y_pred_decision_trees_filtered)
svc_false_positive_rate,svc_true_positive_rate,svc_threshold =_
      ↪roc_curve(y_test,y_pred_svm_filtered)
sns.set_style('whitegrid')
```

```

plt.figure(figsize=(6,4))
plt.title('Receiver Operating Characteristic Curve')


# Plot ROC curves for each classifier with different line styles and colors
plt.plot(lr_false_positive_rate, lr_true_positive_rate, label='Logistic_Regression', color='red', linestyle='-')
plt.plot(nb_false_positive_rate, nb_true_positive_rate, label='Naive Bayes', color='purple', linestyle='-')
plt.plot(rf_false_positive_rate, rf_true_positive_rate, label='Random Forest', color='yellow', linestyle='--')
plt.plot(knn_false_positive_rate, knn_true_positive_rate, label='K-Nearest Neighbor', color='blue', linestyle=':')
plt.plot(dt_false_positive_rate, dt_true_positive_rate, label='Decision Tree', color='orange', linestyle='-')
plt.plot(svc_false_positive_rate, svc_true_positive_rate, label='Support Vector Classifier', color='cyan', linestyle='-')

# Plot the diagonal reference line
plt.plot([0, 1], [0, 1], ls='--', color='.5')

# Add labels and legend
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.legend()

# Show the plot
plt.show()

```



```
[ ]: print("Logistic Regression ROC Curve:")
for fpr, tpr, threshold in zip(lr_false_positive_rate, lr_true_positive_rate,
    ↳lr_threshold):
    print(f"FPR: {fpr}, TPR: {tpr}, Threshold: {threshold}")

print("\nNaive Bayes ROC Curve:")
for fpr, tpr, threshold in zip(nb_false_positive_rate, nb_true_positive_rate,
    ↳nb_threshold):
    print(f"FPR: {fpr}, TPR: {tpr}, Threshold: {threshold}")

print("\nRandom Forest ROC Curve:")
for fpr, tpr, threshold in zip(rf_false_positive_rate, rf_true_positive_rate,
    ↳rf_threshold):
    print(f"FPR: {fpr}, TPR: {tpr}, Threshold: {threshold}")

print("\nK Nearest Neighbors ROC Curve:")
for fpr, tpr, threshold in zip(knn_false_positive_rate, knn_true_positive_rate,
    ↳knn_threshold):
    print(f"FPR: {fpr}, TPR: {tpr}, Threshold: {threshold}")

print("\nDecision Trees ROC Curve:")
for fpr, tpr, threshold in zip(dt_false_positive_rate, dt_true_positive_rate,
    ↳dt_threshold):
```

```

    print(f"FPR: {fpr}, TPR: {tpr}, Threshold: {threshold}")

print("\nSupport Vector Machines ROC Curve:")
for fpr, tpr, threshold in zip(svc_false_positive_rate, svc_true_positive_rate,
    ↪svc_threshold):
    print(f"FPR: {fpr}, TPR: {tpr}, Threshold: {threshold}")

```

Logistic Regression ROC Curve:

FPR: 0.0, TPR: 0.0, Threshold: inf

FPR: 0.2549019607843137, TPR: 0.8737864077669902, Threshold: 1.0

FPR: 1.0, TPR: 1.0, Threshold: 0.0

Naive Bayes ROC Curve:

FPR: 0.0, TPR: 0.0, Threshold: inf

FPR: 0.28431372549019607, TPR: 0.8737864077669902, Threshold: 1.0

FPR: 1.0, TPR: 1.0, Threshold: 0.0

Random Forest ROC Curve:

FPR: 0.0, TPR: 0.0, Threshold: inf

FPR: 0.0196078431372549, TPR: 0.970873786407767, Threshold: 1.0

FPR: 1.0, TPR: 1.0, Threshold: 0.0

K Nearest Neighbors ROC Curve:

FPR: 0.0, TPR: 0.0, Threshold: inf

FPR: 0.0196078431372549, TPR: 1.0, Threshold: 1.0

FPR: 1.0, TPR: 1.0, Threshold: 0.0

Decision Trees ROC Curve:

FPR: 0.0, TPR: 0.0, Threshold: inf

FPR: 0.0196078431372549, TPR: 0.9805825242718447, Threshold: 1.0

FPR: 1.0, TPR: 1.0, Threshold: 0.0

Support Vector Machines ROC Curve:

FPR: 0.0, TPR: 0.0, Threshold: inf

FPR: 0.0392156862745098, TPR: 0.941747572815534, Threshold: 1.0

FPR: 1.0, TPR: 1.0, Threshold: 0.0

```
[ ]: #Ensemble Techniques
```

```

from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import StackingClassifier
from sklearn.pipeline import Pipeline

```

```
[ ]: #Bagging Ensemble (Homogeneous)
```

```

# Create a base classifier (e.g., Decision Tree)
base_classifier = DecisionTreeClassifier(random_state=42)

```

```

# Create a BaggingClassifier
bagging_clf = BaggingClassifier(base_estimator=base_classifier,
    ↪n_estimators=30, random_state=42)

# Train the BaggingClassifier
bagging_clf.fit(X_train, y_train)

# Make predictions
y_pred_bagging = bagging_clf.predict(X_test)

accuracy_bagging = accuracy_score(y_test, y_pred_bagging)
recall_bagging = recall_score(y_test, y_pred_bagging)
precision_bagging = precision_score(y_test, y_pred_bagging)
conf_matrix_bagging = confusion_matrix(y_test, y_pred_bagging)
tn, fp, fn, tp = conf_matrix.ravel()
specificity_after = tn / (tn + fp)
sensitivity_after = tp / (tp + fn)
npv_after = tn / (tn + fn)
f1_score_after = 2 * (precision_bagging * recall_bagging) / (precision_bagging
    ↪+ recall_bagging)
roc_auc = roc_auc_score(y_test, y_pred_bagging)
confusion_matrix_bagging = confusion_matrix(y_test, y_pred_bagging)

# Print evaluation metrics after tuning
print("\nEvaluation Metrics after Bagging:")
print("Accuracy:", accuracy_after)
print("Precision:", precision_after)
print("Recall:", recall_after)
print("Specificity:", specificity_after)
print("F1 Score:", f1_score_after)
print("ROC AUC Score:", roc_auc)
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred_bagging))

plt.figure(figsize=(3, 3))
sns.heatmap(conf_matrix_bagging, annot=True, cmap='YlGnBu', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

```

D:\Anaconda\envs\workingenv\lib\site-packages\sklearn\ensemble_base.py:156:
FutureWarning: `base_estimator` was renamed to `estimator` in version 1.2 and
will be removed in 1.4.
warnings.warn(

Evaluation Metrics after Bagging:

Accuracy: 0.9609756097560975

Precision: 0.9523809523809523

Recall: 0.970873786407767

Specificity: 0.9509803921568627


F1 Score: 0.9852216748768473

ROC AUC Score: 0.9854368932038835

Confusion Matrix:

```
[[102  0]
```

```
 [ 3 100]]
```



```
[ ]: # Ensemble Tuned Models (heterogeneous)
base_models = [
    ('Logistic Regression', grid_search_lr.best_estimator_),
    ('Naive Bayes', naive_bayes_classifier),
    ('Random Forest', grid_search_rf.best_estimator_),
    # ('Decision Tree', grid_search_decision_tree.best_estimator_),
    ('K Nearest Neighbors', grid_search_knn.best_estimator_),
    ('Support Vector Machines', grid_search_svc.best_estimator_)
]

stacking_clf = StackingClassifier(
    estimators=base_models,
    final_estimator=grid_search_decision_tree.best_estimator_
)

# Train the stacking classifier on the training set
```



```

stacking_clf.fit(X_train, y_train)

# Evaluate the stacking classifier on the test set
y_pred_test_stacking = stacking_clf.predict(X_test)
accuracy = accuracy_score(y_test, y_pred_test_stacking)
precision = precision_score(y_test, y_pred_test_stacking)
recall = recall_score(y_test, y_pred_test_stacking)
conf_matrix = confusion_matrix(y_test, y_pred_test_stacking)
tn, fp, fn, tp = conf_matrix.ravel()
specificity = tn / (tn + fp)
sensitivity = tp / (tp + fn)
npv = tn / (tn + fn)
f1_score = 2 * (precision * recall) / (precision + recall)
roc_auc = roc_auc_score(y_test, y_pred_test_stacking)
confusion_matrix_stacking = confusion_matrix(y_test,
y_pred_test_stacking)

# Print evaluation metrics
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("Specificity:", specificity)
print("F1 Score:", f1_score)
print("ROC AUC Score:", roc_auc)

print(confusion_matrix_stacking)


plt.figure(figsize=(3, 3))
sns.heatmap(confusion_matrix_stacking, annot=True, cmap='terrain', fmt='g')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

```

```

Accuracy: 0.975609756097561
Precision: 1.0
Recall: 0.9514563106796117
Specificity: 1.0
F1 Score: 0.9751243781094527
ROC AUC Score: 0.9757281553398058
[[102  0]
 [ 5 98]]

```



```
[ ]: ## Ensemble Tuned Models
## base_models = [
##     ('Logistic Regression', logistic_regression),
##     ('Naive Bayes', naive_bayes_classifier),
##     ('Random Forest', random_forest_classifier),
##     # ('Decision Tree', grid_search_decision_tree.best_estimator_),
##     ('K Nearest Neighbors', knn_classifier),
##     ('Support Vector Machines', svc_classifier)
## ]

## stacking_clf = StackingClassifier(
##     estimators=base_models,
##     final_estimator=decision_tree_classifier
## )

## Train the stacking classifier on the training set
## stacking_clf.fit(X_train, y_train)

## Evaluate the stacking classifier on the test set
## y_pred_test_stacking = stacking_clf.predict(X_test)
## accuracy = accuracy_score(y_test, y_pred_test_stacking)
## precision = precision_score(y_test, y_pred_test_stacking)
## recall = recall_score(y_test, y_pred_test_stacking)
## conf_matrix = confusion_matrix(y_test, y_pred_test_stacking)
## tn, fp, fn, tp = conf_matrix.ravel()
## specificity = tn / (tn + fp)
## sensitivity = tp / (tp + fn)
```

```

# npv = tn / (tn + fn)
# f1_score = 2 * (precision * recall) / (precision + recall)
# roc_auc = roc_auc_score(y_test, y_pred_test_stacking)
# confusion_matrix_stacking = confusion_matrix(y_test,
# y_pred_test_stacking)

# # Print evaluation metrics
# print("Accuracy:", accuracy)
# print("Precision:", precision)
# print("Recall:", recall)
# print("Specificity:", specificity)
# print("F1 Score:", f1_score)
# print("ROC AUC Score:", roc_auc)

# print(confusion_matrix_stacking)

# plt.figure(figsize=(3, 3))
# sns.heatmap(confusion_matrix_stacking, annot=True, cmap='terrain', fmt='g')
# plt.xlabel('Predicted')
# plt.ylabel('True')
# plt.title('Confusion Matrix')
# plt.show()

```

```

[ ]: # # Ensemble Feature Selection
# base_models = [
#     ('Logistic Regression', grid_search_lr.best_estimator_),
#     ('Naive Bayes', naive_bayes_classifier),
#     ('Random Forest', grid_search_rf.best_estimator_),
#     ('Decision Tree', grid_search_decision_tree.best_estimator_),
#     ('K Nearest Neighbors', grid_search_knn.best_estimator_),
#     ('Support Vector Machines', grid_search_svm.best_estimator_)
# ]

# # Define feature selection models
# cv_lr = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
# cv_rf = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)

# feature_selection_models = [
#     ('RFECV Logistic', RFECV(estimator=LogisticRegression(), step=1,
# ↪cv=cv_lr, scoring='accuracy')),
#     ('SelectKBest Naive Bayes', SelectKBest(score_func=mutual_info_classif,
# ↪k=10)),
#     ('SelectKBest KNN', SelectKBest(score_func=mutual_info_classif, k=10)),
#     ('RFECV Random Forest', RFECV(estimator=RandomForestClassifier(), step=1,
# ↪cv=cv_rf, scoring='accuracy')),
#     ('SelectKBest SVM', SelectKBest(score_func=mutual_info_classif, k=11))
# ]

```

```

# estimators = []
# for model, fs_model in feature_selection_models:
#     for base_name, base_model in base_models:
#         estimators.append((f"{model} + {base_name}", Pipeline([('Feature_
↳ Selection', fs_model), (base_name, base_model)])))

# stacking_clf = StackingClassifier(
#     estimators=base_models,
#     final_estimator= grid_search_knn
# )

# # Train the stacking classifier on the training set
# stacking_clf.fit(X_train, y_train)

# # Evaluate the stacking classifier on the test set
# y_pred_test_stacking = stacking_clf.predict(X_test)
# accuracy = accuracy_score(y_test, y_pred_test_stacking)
# precision = precision_score(y_test, y_pred_test_stacking)
# recall = recall_score(y_test, y_pred_test_stacking)
# conf_matrix = confusion_matrix(y_test, y_pred_test_stacking)
# tn, fp, fn, tp = conf_matrix.ravel()
# specificity = tn / (tn + fp)
# sensitivity = tp / (tp + fn)
# npv = tn / (tn + fn)
# f1_score = 2 * (precision * recall) / (precision + recall)
# roc_auc = roc_auc_score(y_test, y_pred_test_stacking)

# # Print evaluation metrics
# print("Accuracy:", accuracy)
# print("Precision:", precision)
# print("Recall:", recall)
# print("Specificity:", specificity)
# print("F1 Score:", f1_score)
# print("ROC AUC Score:", roc_auc)

```

```

[ ]: # # Stacking Ensemble Technique (heterogeneous)
# # Pre-trained and tuned models
# base_models = [
#     ('Logistic Regression', LogisticRegression()),
#     ('Naive Bayes', GaussianNB()),
#     ('Random Forest', RandomForestClassifier()),
#     ('Decision Tree', DecisionTreeClassifier()),
#     ('K Nearest Neighbors', KNeighborsClassifier()),
#     ('Support Vector Machines', SVC())
# ]

```

```

# cv_lr = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
# cv_rf = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)

# # Define feature selection models
# feature_selection_models = [
#     ('RFECV Logistic', RFECV(estimator=LogisticRegression(), step=1,
# ↪cv=cv_lr, scoring='accuracy')),
#     ('SelectKBest Naive Bayes', SelectKBest(score_func=mutual_info_classif,
# ↪k=10)),
#     ('SelectKBest KNN', SelectKBest(score_func=mutual_info_classif, k=10)),
#     ('RFECV Random Forest', RFECV(estimator=RandomForestClassifier(), step=1,
# ↪cv=cv_rf, scoring='accuracy')),
#     ('SelectKBest SVM', SelectKBest(score_func=mutual_info_classif, k=11))
# ]

# estimators = []
# for model, fs_model in feature_selection_models:
#     for base_name, base_model in base_models:
#         estimators.append((f"{model} + {base_name}", Pipeline([('Feature
# ↪Selection', fs_model), (base_name, base_model)])))

# stacking_clf = StackingClassifier(
#     estimators=base_models,
#     final_estimator= decision_trees_rfecv
# )

# # Train the stacking classifier on the training set
# stacking_clf.fit(X_train, y_train)

# # Evaluate the stacking classifier on the test set
# y_pred_test_stacking = stacking_clf.predict(X_test)
# accuracy = accuracy_score(y_test, y_pred_test_stacking)
# precision = precision_score(y_test, y_pred_test_stacking)
# recall = recall_score(y_test, y_pred_test_stacking)
# conf_matrix = confusion_matrix(y_test, y_pred_test_stacking)
# tn, fp, fn, tp = conf_matrix.ravel()
# specificity = tn / (tn + fp)
# sensitivity = tp / (tp + fn)
# npv = tn / (tn + fn)
# f1_score = 2 * (precision * recall) / (precision + recall)
# roc_auc = roc_auc_score(y_test, y_pred_test_stacking)

# # Print evaluation metrics
# print("Accuracy:", accuracy)
# print("Precision:", precision)

```

```
# print("Recall:", recall)
# print("Specificity:", specificity)
# print("F1 Score:", f1_score)
# print("ROC AUC Score:", roc_auc)
```